

Introduction to Code Generation

- We first focus on generating code for a stack machine with accumulator
- We want to run the resulting code on a real machine, e.g., the MIPS/x86 processor (or simulator)
 - The ideas do not change much for abstract machines like LLVM IR abstract machine, except some issues like register-allocation.
- We simulate stack machine instructions using MIPS/x86 instructions and registers.

MIPS architecture

- Prototypical Reduced Instruction Set Computer (RISC)
- Most operations use registers for operands and results
- Use load and store instructions to use values in memory
- 32 general purpose registers (32 bits each)
 - We will use only three registers: \$sp, \$a0, and \$t1 (a temporary register)

x86 architecture

- Complex Instruction Set Computer (CISC)
- Significantly larger opcode set : 400-odd compared to 40-odd in RISC
- Opcodes often can operate on both registers and/or memory
 - Do not necessarily need separate load/store instructions
- 8 general purpose registers (32 bits each)
 - We will use %esp, %eax, and %ecx (a temporary register)
 - Intel engineers felt that it is better to provide more opcodes and less registers. Use on-chip real-estate for more functional units and logic (by saving space through a shorter register file and its connections).

Need some code-generation invariants

- The accumulator is kept in MIPS register \$a0 (or x86 register %eax)
- The stack is kept in memory
 - Grows downwards towards lower addresses
 - Standard convention on both MIPS and x86
- For MIPS
 - The next location of the stack is kept in MIPS register \$sp. The top of the stack is at address \$sp+4
- For x86
 - The top of the stack is at address %esp. The next location of the stack is at address %esp-4.
- The name "sp" stands for stack-pointer.

MIPS opcodes (relevant only)

- lw reg1, offset(reg2)
 - Load 32-bit word from address reg2+offset into reg1
- add reg1, reg2, reg3
 - reg1 <-- reg2+reg3
- sw reg1, offset(reg2)
 - Store 32-bit word reg1 to address reg2+offset
- addiu reg1, reg2, imm
 - reg1 <-- reg2+imm
 - "u" means overflow is not checked (overflow means different things for signed and unsigned interpretation of the registers)
- li reg, imm
 - reg <-- imm

x86 opcodes (relevant only)

- movl %reg1/(memaddr1)/\$imm, %reg2/(memaddr2)
 - Move 32-bit word from register reg1 (or address memaddr1 or the immediate value itself) into reg2 or to memory address memaddr2

- Captures several opcodes in one mnemonic (load, store, li, move-register, etc.). More powerful than RISC, e.g., MIPS cannot move immediate value directly to memory
- `add %reg1/(memaddr1)/$imm, %reg2/(memaddr2)`
 - `%reg2/(memaddr2) <-- reg1/(memaddr1)/imm + %reg2/(memaddr2)`
 - Overflow is always computed for both signed/unsigned arithmetic. Happens in parallel so not in critical performance path, but switches more transistors (more power)
- `push %reg/(memaddr)/$imm`
 - `(%esp-4) <-- reg/(memaddr)/imm; %esp <-- %esp-4`
- `pop %reg/(memaddr)/$imm`
 - `reg/(memaddr)/imm <-- (%esp); %esp <-- %esp+4`
- push/pop are "higher-level" opcodes: enables faster execution paths for these common operations

The stack-machine code for 7+5 in MIPS:

```
acc <-- 7          : li $a0, 7
push acc          : sw $a0, 0($sp)
                  : addiu $sp, $sp, -4
acc <-- 5          : li $a0, 5
acc < acc + top_of_stack : lw $t1, 4($sp)
                  : add $a0, $a0, $t1
pop               : addiu $sp, $sp, 4
```

The stack-machine code for 7+5 in x86:

```
acc <-- 7          : movl $7,%eax
push acc          : pushl %eax
acc <-- 5          : movl $5, %eax
acc < acc + top_of_stack : addl (%esp),%eax
pop               : popl %ecx      #just pop the register to unused register ecx
```

A more optimized version was possible in x86:

```
acc <-- 7          : push $7
push acc          :
acc <-- 5          : mov $5, %eax
acc < acc + top_of_stack : add (%esp), %eax
pop               : pop %ecx      #just pop the register to unused register ecx
```