

Optimization Overview

Most complexity in modern compilers is in the optimizer

- Also by far, the largest phase in terms of compile-time and in terms of source code size (recall: lexing, parsing, semantic analysis, optimization, code generation)

When should we perform optimizations?

- On AST?
 - Pro: Machine independent
 - Con: Too high level
 - Too abstract --- need to have more details to be able to express the "kind" of machine for which the AST needs to be compiled (e.g., register or stack machine or quantum computer)
- On assembly language
 - Pro: exposes optimization opportunities
 - Con: may be too low level, making optimization difficult (need to undo/redo certain decisions)
 - Con: machine dependent
 - Con: must reimplement optimizations when retargeting
- On IR
 - Pro: can be machine independent, if designed well (can represent a large family of machines)
 - Pro: can expose optimization opportunities, if designed well.
 - Con: IR design critical for exposing optimization opportunities.

We will be looking at optimizations on an IR which has the following grammar:

```
P --> S P | S
S --> id := id op id
      | id := op id
      | id := id
      | push id
      | id := pop
      | if id relop id goto L
      | L:
      | jump L
```

- id's are register names
- constants can replace ids
- typical operators: +, -, *

A basic block is a maximal sequence of instructions with

- no labels (except at the first instruction), and
- no jumps (except in the last instruction)

Idea:

- Cannot jump into a basic block (except at beginning)
- Cannot jump out of a basic block (except at end)
- A basic block is a single-entry, single-exit, straight-line code segment

Once we reach the start of a basic block, we are guaranteed to execute all instructions in the BB. Furthermore, the only way into the basic block is through the first statement.

Consider the basic block:

```
1. L:
2.   t := 2*x
3.   w := t + x
4.   if w > 0 goto L'
```

(3) executes only after (2)

- We can change (3) to $w := 3 * x$
- Can we eliminate (2) as well? Need to be sure that t is not used in other basic blocks.

A control-flow graph is a directed graph with

- Basic blocks as nodes
- An edge from block A to block B if the execution can pass from the last instruction in A to the first instruction in B
 - e.g., the last instruction in A is `jump Lb`
 - e.g., the last instruction in A is `if id relop id then goto Lb`
 - e.g., execution can fall-through from block A to block B

Example control-flow graph:

BB1:

```
x := 1
i := 1
```

BB1 --> BB2

BB2:

L:

```
x := x * x
i := i + 1
if i < 10 goto L
```

BB2 --> BB2

BB2 --> BB3

The body of a method (or procedure) can be represented as a control-flow graph. There is one initial node (entry node). All "return" nodes are terminal.

Optimization seeks to improve a program's resource utilization

- Execution time (most often)
- Code size
- Memory usage
- Network messages sent, disk accesses, etc.
- Power consumption

Optimization should not alter what the program computes

- The answer must still be the same.

For languages like C, there are typically three granularities of optimization

- Local optimizations
 - Apply to a basic block in isolation
- Global optimizations
 - Apply to a control-flow graph (method body) in isolation
- Inter-procedural optimizations
 - Apply across method boundaries

Production compilers do all these types of optimizations. In general, easiest to implement local optimizations and hardest to implement inter-procedural optimizations.

Criteria for evaluating an optimization

- Payoff. What is a good payoff?

- Frequency of occurrence
- Ease of Implementation
- Compilation Time

In practice, often a conscious decision is made not to implement the fanciest optimization known. Why?

- Some optimizations are hard to implement
- Some optimizations are costly in compilation time
- Some optimizations have low payoff. But hard to establish the payoff. Often one optimization may trigger another one, and this is hard to predict in advance.
- Many fancy optimizations are all three!

Current state-of-the-art: the goal is "Maximum benefit for minimum cost"